

Assignment 2: Lending Club Data

Huey Ning Lok

February 12, 2019

The Lending Club is a platform which allows the crowdfunding of various loans. Various investors are able to browse the profiles of people applying for loans and decide whether or not to help fund them. The Lending club has made an anonymized set of data available for anyone to study at: <https://www.lendingclub.com/info/download-data.action>

In this assignment, I will build a model that predicts the largest loan amount that will be successfully funded for given individual. This model can then be used to advise applicants on how much they could apply for.

1 Data Cleaning and Preprocessing

1.1 Importing Libraries and Importing Data

The data imported is the 2015 Approved Loan and 2015 Declined Loan csv data from the lending club website.

```
In [ ]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [202]: #import loan data
loan = pd.read_csv("LoanStats3d.csv", skiprows=1)
#column names
print(list(loan))
#peak at the df
loan.head()
```

```
/anaconda3/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2785:
DtypeWarning: Columns (0,19,55) have mixed types. Specify dtype option on import or
set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)
```

```
['id', 'member_id', 'loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'term', 'int_rate',
'installment', 'grade', 'sub_grade', 'emp_title', 'emp_length', 'home_ownership',
'annual_inc', 'verification_status', 'issue_d', 'loan_status', 'pymnt_plan', 'url',
'desc', 'purpose', 'title', 'zip_code', 'addr_state', 'dti', 'delinq_2yrs',
'earliest_cr_line', 'inq_last_6mths', 'mths_since_last_delinq',
'mths_since_last_record', 'open_acc', 'out_prncp', 'revol_bal', 'revol_util',
'total_acc', 'initial_list_status', 'out_prncp_inv', 'total_pymnt',
'total_pymnt_inv', 'total_rec_prncp', 'total_rec_int', 'total_rec_late_fee',
'recoveries', 'collection_recovery_fee', 'last_pymnt_d', 'last_pymnt_amnt',
```

```

'next_pymnt_d', 'last_credit_pull_d', 'collections_12_mths_ex_med',
'mths_since_last_major_derog', 'policy_code', 'application_type', 'annual_inc_joint',
'dti_joint', 'verification_status_joint', 'acc_now_delinq', 'tot_coll_amt',
'tot_cur_bal', 'open_acc_6m', 'open_act_il', 'open_il_12m', 'open_il_24m',
'mths_since_rcnt_il', 'total_bal_il', 'il_util', 'open_rv_12m', 'open_rv_24m',
'max_bal_bc', 'all_util', 'total_rev_hi_lim', 'inq_fi', 'total_cu_tl', 'inq_last_12m',
'acc_open_past_24mths', 'avg_cur_bal', 'bc_open_to_buy', 'bc_util',
'chargeoff_within_12_mths', 'delinq_amnt', 'mo_sin_old_il_acct',
'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op', 'mo_sin_rcnt_tl', 'mort_acc',
'mths_since_recent_bc', 'mths_since_recent_bc_dlq', 'mths_since_recent_inq',
'mths_since_recent_revol_delinq', 'num_accts_ever_120_pd', 'num_actv_bc_tl',
'num_actv_rev_tl', 'num_bc_sats', 'num_bc_tl', 'num_il_tl', 'num_op_rev_tl',
'num_rev_accts', 'num_rev_tl_bal_gt_0', 'num_sats', 'num_tl_120dpd_2m',
'num_tl_30dpd', 'num_tl_90g_dpd_24m', 'num_tl_op_past_12m', 'pct_tl_nvr_dlq',
'percent_bc_gt_75', 'pub_rec_bankruptcies', 'tax_liens', 'tot_hi_cred_lim',
'total_bal_ex_mort', 'total_bc_limit', 'total_il_high_credit_limit',
'revol_bal_joint', 'sec_app_earliest_cr_line', 'sec_app_inq_last_6mths',
'sec_app_mort_acc', 'sec_app_open_acc', 'sec_app_revol_util', 'sec_app_open_act_il',
'sec_app_num_rev_accts', 'sec_app_chargeoff_within_12_mths',
'sec_app_collections_12_mths_ex_med', 'sec_app_mths_since_last_major_derog',
'hardship_flag', 'hardship_type', 'hardship_reason', 'hardship_status',
'deferral_term', 'hardship_amount', 'hardship_start_date', 'hardship_end_date',
'payment_plan_start_date', 'hardship_length', 'hardship_dpd', 'hardship_loan_status',
'orig_projected_additional_accrued_interest', 'hardship_payoff_balance_amount',
'hardship_last_payment_amount', 'disbursement_method', 'debt_settlement_flag',
'debt_settlement_flag_date', 'settlement_status', 'settlement_date',
'settlement_amount', 'settlement_percentage', 'settlement_term']

```

```

Out[202]:      id  member_id  loan_amnt  funded_amnt  funded_amnt_inv      term \
0   NaN         NaN    24700.0    24700.0         24700.0    36 months
1   NaN         NaN    17925.0    17925.0         17925.0    60 months
2   NaN         NaN     9000.0     9000.0          9000.0    36 months
3   NaN         NaN    11550.0    11550.0         11550.0    60 months
4   NaN         NaN    12000.0    12000.0         12000.0    60 months

      int_rate  installment  grade  sub_grade      ...      \
0    11.99%         820.28      C         C1      ...      \
1    17.27%         448.09      D         D3      ...      \
2     8.49%         284.07      B         B1      ...      \
3    16.59%         284.51      D         D2      ...      \
4     9.80%         253.79      B         B3      ...      \

      hardship_payoff_balance_amount  hardship_last_payment_amount      \
0                                NaN                             NaN
1                                NaN                             NaN
2                                NaN                             NaN
3                                NaN                             NaN
4                                NaN                             NaN

      disbursement_method  debt_settlement_flag  debt_settlement_flag_date      \
0                   Cash                        N                        NaN
1                   Cash                        N                        NaN

```

2	Cash	N	NaN
3	Cash	N	NaN
4	Cash	N	NaN

	settlement_status	settlement_date	settlement_amount	settlement_percentage \
0	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

	settlement_term
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 145 columns]

```
In [203]: #import reject data
reject = pd.read_csv("RejectStatsD.csv",skiprows=1)
#column names
print(list(reject))
#peak at the df
reject.head()
```

```
['Amount Requested', 'Application Date', 'Loan Title', 'Risk_Score', 'Debt-To-Income
Ratio', 'Zip Code', 'State', 'Employment Length', 'Policy Code']
```

```
Out[203]:
```

	Amount Requested	Application Date	Loan Title	Risk_Score \
0	30000.0	2015-01-01	debt_consolidation	681.0
1	5000.0	2015-01-01	debt_consolidation	648.0
2	10000.0	2015-01-01	Debt consolidation	721.0
3	10000.0	2015-01-01	major_purchase	659.0
4	5000.0	2015-01-01	debt_consolidation	501.0

	Debt-To-Income Ratio	Zip Code	State	Employment Length	Policy Code
0	35.65%	958xx	CA	< 1 year	0
1	10.62%	945xx	CA	< 1 year	0
2	10.02%	750xx	TX	7 years	0
3	19.05%	853xx	AZ	< 1 year	0
4	10.73%	475xx	IN	< 1 year	0

1.2 Choosing which columns to keep

I can only keep columns that exist in both the approved and rejected loan datasets. The table below shows the column names that exist in both the loan data and reject data that share the same definition (as defined by the data dictionary on the lending data website).

Loan Data	Reject Data	Definition
loan_amnt	Amount Requested	The total amount requested by the borrower
title	Loan Title	The loan title provided by the borrower
dti	Debt-To-Income Ratio	A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LC loan, divided by the borrower's self-reported monthly income.
zip_code	Zip Code	The first 3 numbers of the zip code provided by the borrower in the loan application.
addr_state	State	The state provided by the borrower in the loan application
emp_length	Employment Length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
policy_code	Policy Code	publicly available policy_code=1; new products not publicly available policy_code=2

In [204]: *#Drop all columns that do not exist in both dataframes.*

```
loan =
loan[['loan_amnt', 'title', 'dti', 'zip_code', 'addr_state', 'emp_length', 'policy_code']]
loan['Status'] = 1
print(loan.shape)
loan.head()
```

(421097, 8)

```
Out[204]:
```

	loan_amnt	title	dti	zip_code	addr_state	emp_length	\
0	24700.0	Business	16.06	577xx	SD	10+ years	
1	17925.0	Debt consolidation	27.78	432xx	OH	10+ years	
2	9000.0	Debt consolidation	8.43	346xx	FL	8 years	
3	11550.0	Credit card refinancing	21.07	436xx	OH	5 years	
4	12000.0	Debt consolidation	23.84	660xx	KS	10+ years	

	policy_code	Status
0	1.0	1
1	1.0	1
2	1.0	1
3	1.0	1
4	1.0	1

```
In [205]: reject = reject[['Amount Requested', 'Loan Title', 'Debt-To-Income Ratio', 'Zip
Code', 'State', 'Employment Length', 'Policy Code']]
reject['Status'] = 0
print(reject.shape)
reject.head()
```

(2859379, 8)

```
Out[205]:
```

	Amount Requested	Loan Title	Debt-To-Income Ratio	Zip Code	State	\
0	30000.0	debt_consolidation	35.65%	958xx	CA	

1	5000.0	debt_consolidation	10.62%	945xx	CA
2	10000.0	Debt consolidation	10.02%	750xx	TX
3	10000.0	major_purchase	19.05%	853xx	AZ
4	5000.0	debt_consolidation	10.73%	475xx	IN

	Employment Length	Policy Code	Status
0	< 1 year	0	0
1	< 1 year	0	0
2	7 years	0	0
3	< 1 year	0	0
4	< 1 year	0	0

```
In [206]: #rename columns in loan data so that it is the same as reject data
loan.columns = list(reject)
data = pd.concat([loan,reject])
print(data.shape)
data.head()
```

(3280476, 8)

```
Out[206]:
```

	Amount Requested	Loan Title	Debt-To-Income Ratio	Zip Code	\
0	24700.0	Business	16.06	577xx	
1	17925.0	Debt consolidation	27.78	432xx	
2	9000.0	Debt consolidation	8.43	346xx	
3	11550.0	Credit card refinancing	21.07	436xx	
4	12000.0	Debt consolidation	23.84	660xx	

	State	Employment Length	Policy Code	Status
0	SD	10+ years	1.0	1
1	OH	10+ years	1.0	1
2	FL	8 years	1.0	1
3	OH	5 years	1.0	1
4	KS	10+ years	1.0	1

```
In [207]: data.info(verbose=True, null_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3280476 entries, 0 to 2859378
Data columns (total 8 columns):
Amount Requested      3280474 non-null float64
Loan Title            3280342 non-null object
Debt-To-Income Ratio  3280472 non-null object
Zip Code              3280474 non-null object
State                 3280474 non-null object
Employment Length     3182914 non-null object
Policy Code           3280474 non-null float64
Status                3280476 non-null int64
dtypes: float64(2), int64(1), object(5)
memory usage: 225.3+ MB
```

```
In [208]: #drop NaN values from data
data.dropna(inplace=True)
data.info(verbose=True, null_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3182789 entries, 0 to 2859378
Data columns (total 8 columns):
Amount Requested      3182789 non-null float64
Loan Title            3182789 non-null object
Debt-To-Income Ratio  3182789 non-null object
Zip Code              3182789 non-null object
State                 3182789 non-null object
Employment Length     3182789 non-null object
Policy Code           3182789 non-null float64
Status                3182789 non-null int64
dtypes: float64(2), int64(1), object(5)
memory usage: 218.5+ MB
```

```
In [209]: #convert debt-to-income ratio and employment length to numeric data types
data['Debt-To-Income Ratio'] = data['Debt-To-Income
Ratio'].astype(str).str.extract('(\d+)').astype(float)
data['Employment Length'] = data['Employment Length'].str.extract('(\d+)').astype(int)

print(data.info(verbose=True, null_counts=True))
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3182789 entries, 0 to 2859378
Data columns (total 8 columns):
Amount Requested      3182789 non-null float64
Loan Title            3182789 non-null object
Debt-To-Income Ratio  3182789 non-null float64
Zip Code              3182789 non-null object
State                 3182789 non-null object
Employment Length     3182789 non-null int64
Policy Code           3182789 non-null float64
Status                3182789 non-null int64
dtypes: float64(3), int64(2), object(3)
memory usage: 218.5+ MB
None
```

```
Out[209]:
```

	Amount Requested	Loan Title	Debt-To-Income Ratio	Zip Code	\
0	24700.0	Business	16.0	577xx	
1	17925.0	Debt consolidation	27.0	432xx	
2	9000.0	Debt consolidation	8.0	346xx	
3	11550.0	Credit card refinancing	21.0	436xx	
4	12000.0	Debt consolidation	23.0	660xx	

	State	Employment Length	Policy Code	Status
0	SD	10	1.0	1
1	OH	10	1.0	1
2	FL	8	1.0	1
3	OH	5	1.0	1
4	KS	10	1.0	1

1.3 Closer Analysis for In-depth Preprocessing

```
In [210]: #basic stats for numerical columns
data.describe(include = [np.number])
```

```
Out [210]:
```

	Amount Requested	Debt-To-Income Ratio	Employment Length \
count	3.182789e+06	3.182789e+06	3.182789e+06
mean	1.435574e+04	1.071487e+02	2.113442e+00
std	1.872363e+04	2.039714e+03	2.519112e+00
min	1.000000e+03	0.000000e+00	1.000000e+00
25%	5.000000e+03	7.000000e+00	1.000000e+00
50%	1.000000e+04	1.800000e+01	1.000000e+00
75%	2.000000e+04	3.200000e+01	1.000000e+00
max	4.500000e+05	1.019525e+06	1.000000e+01

	Policy Code	Status
count	3.182789e+06	3.182789e+06
mean	1.302741e-01	1.247814e-01
std	3.446672e-01	3.304710e-01
min	0.000000e+00	0.000000e+00
25%	0.000000e+00	0.000000e+00
50%	0.000000e+00	0.000000e+00
75%	0.000000e+00	0.000000e+00
max	2.000000e+00	1.000000e+00

```
In [211]: #basic stats for categorical columns
data.describe(include = ['O'])
```

```
Out [211]:
```

	Loan Title	Zip Code	State
count	3182789	3182789	3182789
unique	55	998	49
top	debt_consolidation	112xx	CA
freq	1257311	33264	400190

1.4 Observations

- The Policy Code column appears to contain zeros, although the data dictionary only includes definitions for policy codes of ones and twos.
 - Since more than half the dataset contains policy code values of 0, and there is no way to intuitively predict what policy code == 0 could stand for, I will drop the policy code column.
- There are 55 unique Loan Title values, 998 unique Zip Code values, and 49 unique State values.
 - Due to the high amount of unique Zip Code values, I will drop the Zip Code column as it may be quite messy to one-hot-encode 998 unique values, and Zip Code information can also be indirectly inferred from State values.
 - The Loan Title values will be analyzed and manually sorted into general, representative categories.
 - The State values will be kept as they are.
- From the summary stats tables above, it looks like there is a huge outlier in the Debt-To-Income Ratio category - the max value is 10^5 times bigger than the 75% percentile.¹
 - While the outlier is erroneous, it is also feasible; there could be an individual out there with very high debt and very low income. Excluding the outlier could then lead to a biased dataset. On the other hand, the outlier may also be a result of human error - in which case the resulting model would not be representative of the data. I have decided to keep the outlier since no robust tests were conducted to determine whether it can be considered a legitimate outlier, and it is possible for an individual to have a very high debt-to-income ratio.

1.4.1 Drop Policy Code Column

```
In [212]: #Checking the number of rows with policy code values == 0
print(f"Number of rows where policy code == 0: {sum(data['Policy Code'] == 0)}")

#drop Policy Code column
data.drop(columns=['Policy Code'], inplace=True)

#peak at the df
data.head()
```

Number of rows where policy code == 0: 2776895

```
Out [212]:
```

	Amount Requested	Loan Title	Debt-To-Income Ratio	Zip Code	\
0	24700.0	Business	16.0	577xx	
1	17925.0	Debt consolidation	27.0	432xx	
2	9000.0	Debt consolidation	8.0	346xx	
3	11550.0	Credit card refinancing	21.0	436xx	

¹**#descriptivestats:** By observing a summary of the descriptive stats of the dataframe, I was able to manually detect irregular discrepancies in the data. In this case, the policy code contains 0 as a minimum (which is not included in the data dictionary), and the Debt-To-Income Ratio has a stark outlier - this is inferred by observing the max value ($1.02e+06$), which is a few orders of magnitude bigger than the 75% percentile ($3.2e+01$). Although the descriptive stats isn't a robust method of analyzing data irregularities, any stark irregularities can be intuitively singled out.

4	12000.0	Debt consolidation	23.0	660xx
---	---------	--------------------	------	-------

	State	Employment Length	Status
0	SD	10	1
1	OH	10	1
2	FL	8	1
3	OH	5	1
4	KS	10	1

1.4.2 Drop Zip Code Column

```
In [213]: #Drop the Zip code column
data.drop(columns=['Zip Code'], inplace=True)

#peak at the df
data.head()
```

```
Out[213]:
```

	Amount Requested	Loan Title	Debt-To-Income Ratio	State	\
0	24700.0	Business	16.0	SD	
1	17925.0	Debt consolidation	27.0	OH	
2	9000.0	Debt consolidation	8.0	FL	
3	11550.0	Credit card refinancing	21.0	OH	
4	12000.0	Debt consolidation	23.0	KS	

	Employment Length	Status
0	10	1
1	10	1
2	8	1
3	5	1
4	10	1

1.4.3 Sort and One-Hot-Encode Loan Title Column

```
In [214]: #Get value counts of loan title column
loan_titles_value_counts = data['Loan Title'].value_counts()
print(loan_titles_value_counts)
```

debt_consolidation	1257311
other	339305
Debt consolidation	315246
credit_card	299620
home_improvement	145450
Credit card refinancing	128477
car	108601
moving	86701
Business Loan	83657
medical	80339
major_purchase	80119
small_business	66971
house	38949
vacation	35370
Home improvement	34953
Other	27938
Major purchase	11378

Business	11082
renewable_energy	6949
Medical expenses	5831
Car financing	5675
Moving and relocation	3882
Vacation	3457
Business Line Of Credit	2536
Home buying	2485
Green loan	478
Credit Card/Auto Repair	1
freeup	1
Pay off Lowes Card	1
SAVE	1
thad31	1
new kitchen for momma!	1
Prescription Drug and Medical Costs	1
althea9621	1
Student Loan	1
dougie03	1
10 months away from being an RN	1
New Baby and New House (CC Consolidate)	1
Business Advertising Loan	1
loan	1
considerate	1
Trying to come back to reality!	1
Simple Loan Until Contract Is Completed	1
Consolidation Loan	1
Need a decent rate on car financing	1
educational	1
Consolidate debt	1
Small Business Expansion	1
DebtC	1
Auto Financing	1
odymeds	1
new day	1
Paying off higher interest cards & auto	1
Learning and training	1
smmoore2	1
Name: Loan Title, dtype: int64	

From the value count above, we can see the most frequently used loan titles (any loan title with value > 1). There are some duplicates among these loan titles (e.g. debt_consolidation & Debt consolidation; home_improvement & Home Improvement), and titles which belong to the same general category (e.g. renewable_energy & Green loan; Home improvement & Home Buying).

General categories will be created and the duplicates as well as titles with similar meaning will be sorted into the categories as follows:

General Category	Loan Titles Included
Debt Consolidation	debt_consolidation, Debt consolidation,
Credit Card	credit_card, Credit card refinancing
House	Home improvement, home_improvement, house, Home buying
Car	car, Car financing

General Category	Loan Titles Included
Major Purchase	major_purchase, Major purchase
Moving	Moving and relocation, moving
Business	Business Loan, Business, small_business, Business Line Of Credit
Medical	medical, Medical expenses
Vacation	vacation, Vacation
Green Loan	renewable_energy, Green loan
Other	other, Other, all loan titles with value count == 1

```
In [215]: rare_loan_titles = data['Loan
Title'].isin(loan_titles_value_counts.index[loan_titles_value_counts == 1])
print(f"There are {sum(rare_loan_titles)} loan titles with a value count of 1.\n")

#change value_count == 1 loan titles to "Other"
data.loc[rare_loan_titles, 'Loan Title'] = "Other"

print(data['Loan Title'].value_counts())
```

There are 29 loan titles with a value count of 1.

```
debt_consolidation    1257311
other                  339305
Debt consolidation    315246
credit_card           299620
home_improvement      145450
Credit card refinancing 128477
car                   108601
moving                 86701
Business Loan         83657
medical               80339
major_purchase        80119
small_business        66971
house                 38949
vacation              35370
Home improvement      34953
Other                 27967
Major purchase        11378
Business              11082
renewable_energy      6949
Medical expenses      5831
Car financing         5675
Moving and relocation  3882
Vacation              3457
Business Line Of Credit 2536
Home buying           2485
Green loan            478
Name: Loan Title, dtype: int64
```

```
In [217]: #Sort loan titles as according to the General Category dictionary above
```

```
#loan titles dictionary
loan_titles_dct = {
    'Debt consolidation': 'Debt Consolidation',
    'debt_consolidation': 'Debt Consolidation',
    'credit_card': 'Credit Card',
    'Credit card refinancing': 'Credit Card',
    'Home improvement': 'House',
```

```

        'Home improvement': 'House',
        'home_improvement': 'House',
        'house': 'House',
        'Home buying': 'House',
        'car': 'Car',
        'Car financing': 'Car',
        'major_purchase': 'Major Purchase',
        'Major purchase': 'Major Purchase',
        'Moving and relocation': 'Moving',
        'moving': 'Moving',
        'Business Loan': 'Business',
        'small_business': 'Business',
        'Business': 'Business',
        'Business Line Of Credit': 'Business',
        'medical': 'Medical',
        'Medical expenses': 'Medical',
        'vacation': 'Vacation',
        'Vacation': 'Vacation',
        'renewable_energy': 'Green Loan',
        'Green loan': 'Green Loan',
        'other': 'Other',
        'Other': 'Other'
    }

```

```
data['Loan Title'] = data['Loan Title'].map(loan_titles_dct)
```

```
print(data['Loan Title'].value_counts())
```

```

Debt Consolidation    1572557
Credit Card          428097
Other                 367272
House                221837
Business             164246
Car                  114276
Major Purchase        91497
Moving               90583
Medical              86170
Vacation             38827
Green Loan           7427
Name: Loan Title, dtype: int64

```

```
In [222]: #Check for NaN values
```

```
print(f"Presence of NaN values: {data['Loan Title'].isnull().values.any()}")
```

```
Presence of NaN values: False
```

```
In [236]: #One-hot-encode the Loan Title column
```

```
from sklearn import preprocessing
```

```
mlb = preprocessing.MultiLabelBinarizer()
```

```

loan_titles_df = pd.DataFrame(mlb.fit_transform(data['Loan Title'].str.split(',
')), columns=mlb.classes_, index=data.index)
print(loan_titles_df.shape)
loan_titles_df.head()

```

```
(3182789, 11)
```

```

Out[236]:   Business  Car  Credit Card  Debt Consolidation  Green Loan  House  \
          0         1         0             0             0         0

```

1	0	0	0	1	0	0
2	0	0	0	1	0	0
3	0	0	1	0	0	0
4	0	0	0	1	0	0

	Major Purchase	Medical	Moving	Other	Vacation
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

1.4.4 One-Hot-Encode State Column

```
In [232]: #Observe and analyze State values
          print(data['State'].value_counts())
```

```
CA    400190
TX    283460
NY    246080
FL    241847
GA    117447
IL    116277
PA    115775
OH    113283
NJ    99495
NC    97956
VA    88273
MI    85254
MD    68073
TN    66279
AZ    66238
MA    64275
WA    58178
IN    57982
CO    56771
AL    55329
MO    55078
SC    50255
LA    49288
MN    43130
WI    42167
CT    39999
NV    39596
KY    39349
OK    35139
OR    32533
AR    31347
MS    29836
KS    26914
UT    21207
NM    18750
HI    18515
WV    17398
NH    13273
RI    12846
```

```

DE      10471
NE      8723
MT      8206
AK      7313
SD      6743
DC      6595
VT      6448
WY      6252
ME      3689
ND      3267
Name: State, dtype: int64

```

In [238]: *#One-hot-encode the State column*

```

mlb = preprocessing.MultiLabelBinarizer()

state_df = pd.DataFrame(mlb.fit_transform(data['State'].str.split(',
')),columns=mlb.classes_, index=data.index)
print(state_df.shape)
state_df.head()

```

(3182789, 49)

```

Out [238]:
      AK  AL  AR  AZ  CA  CO  CT  DC  DE  FL  ...  SD  TN  TX  UT  VA  VT  WA  WI  \
0      0   0   0   0   0   0   0   0   0   0  ...   1   0   0   0   0   0   0   0
1      0   0   0   0   0   0   0   0   0   0  ...   0   0   0   0   0   0   0   0
2      0   0   0   0   0   0   0   0   0   1  ...   0   0   0   0   0   0   0   0
3      0   0   0   0   0   0   0   0   0   0  ...   0   0   0   0   0   0   0   0
4      0   0   0   0   0   0   0   0   0   0  ...   0   0   0   0   0   0   0   0

      WV  WY
0      0   0
1      0   0
2      0   0
3      0   0
4      0   0

```

[5 rows x 49 columns]

1.5 Create the Full Dataset

In [265]: *#combining everything back together:*

```

full_data = pd.concat([data,loan_titles_df,state_df], axis = 1)
full_data.head()

```

```

Out [265]:
      Amount Requested      Loan Title  Debt-To-Income Ratio State \
0          24700.0      Business          16.0      SD
1          17925.0  Debt Consolidation          27.0      OH
2           9000.0  Debt Consolidation           8.0      FL
3          11550.0      Credit Card          21.0      OH
4          12000.0  Debt Consolidation          23.0      KS

      Employment Length  Status  Business  Car  Credit Card  Debt Consolidation \

```

0		10	1	1	0	0	0
1		10	1	0	0	0	1
2		8	1	0	0	0	1
3		5	1	0	0	1	0
4		10	1	0	0	0	1

	...	SD	TN	TX	UT	VA	VT	WA	WI	WV	WY
0	...	1	0	0	0	0	0	0	0	0	0
1	...	0	0	0	0	0	0	0	0	0	0
2	...	0	0	0	0	0	0	0	0	0	0
3	...	0	0	0	0	0	0	0	0	0	0
4	...	0	0	0	0	0	0	0	0	0	0

[5 rows x 66 columns]

```
In [361]: full_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3182789 entries, 0 to 2859378
Data columns (total 66 columns):
Amount Requested      float64
Loan Title             object
Debt-To-Income Ratio  float64
State                 object
Employment Length     int64
Status                int64
Business              int64
Car                   int64
Credit Card           int64
Debt Consolidation    int64
Green Loan            int64
House                 int64
Major Purchase        int64
Medical               int64
Moving                int64
Other                 int64
Vacation              int64
AK                    int64
AL                    int64
AR                    int64
AZ                    int64
CA                    int64
CO                    int64
CT                    int64
DC                    int64
DE                    int64
FL                    int64
GA                    int64
HI                    int64
IL                    int64
IN                    int64
KS                    int64
KY                    int64
LA                    int64
MA                    int64
MD                    int64
```

ME	int64
MI	int64
MN	int64
MO	int64
MS	int64
MT	int64
NC	int64
ND	int64
NE	int64
NH	int64
NJ	int64
NM	int64
NV	int64
NY	int64
OH	int64
OK	int64
OR	int64
PA	int64
RI	int64
SC	int64
SD	int64
TN	int64
TX	int64
UT	int64
VA	int64
VT	int64
WA	int64
WI	int64
WV	int64
WY	int64

dtypes: float64(2), int64(62), object(2)
memory usage: 1.6+ GB

2 Part 2: Modeling & Evaluation

Since the goal is to predict the largest loan amount that will be successfully funded for a given individual, I will be training a logistic regression model to predict the probability that a given individual would have their requested amount approved based on the following independent variables:

- Amount Requested
- Debt-To-Income Ratio
- Employment Length
- Loan Title (one-hot-encoded)
- State (one-hot-encoded)

The response variables will be the status of the loan:

- Status
 - 1: Approved
 - 0: Rejected

2.1 Balancing the Response Class

The response class (Status of Loan) has a rather unbalanced approval to rejection class ratio (~0.14). To balance the ratio of approval to rejection (so that the model is properly trained on predicting either class), I will randomly sample from the rejection set in order to cut it down to the length of the approval set. Lastly, I combine both sets back together.

```
In [372]: from sklearn.model_selection import train_test_split

#check ratio of approval/rejection class in full_data
full_data_approvals = full_data[full_data['Status']==1]
full_data_rejections = full_data[full_data['Status']==0]

print(f'Number of approvals in original data: {len(full_data_approvals)}')
print(f'Number of rejections in original data: {len(full_data_rejections)}')
print(f'Ratio of approval/rejection:
{len(full_data_approvals)/len(full_data_rejections)}\n')

#remove some rejection data to obtain a dataset with balanced predictor classes
full_data_rejection_sample = full_data_rejections.sample(n = len(full_data_approvals))
full_data_balanced = pd.concat([full_data_rejection_sample,full_data_approvals])

print(f'Number of approvals in balanced data:
{len(full_data_balanced[full_data_balanced["Status"]==1])}')
print(f'Number of rejections in balanced data:
{len(full_data_balanced[full_data_balanced["Status"]==0])}')
print(f'Ratio of approval/rejection: {len(full_data_balanced[full_data_balanced["Status"
]==1])/len(full_data_balanced[full_data_balanced["Status"]==0])}')
```

```
Number of approvals in original data: 397153
Number of rejections in original data: 2785636
Ratio of approval/rejection: 0.1425717502214934
```

```
Number of approvals in balanced data: 397153
Number of rejections in balanced data: 397153
Ratio of approval/rejection: 1.0
```

2.2 Preventing Overfit: Splitting Data & Cross-Validation

The balanced dataset is split into training and test sets to prevent overfitting. Cross-validation will be conducted on the training set to finetune the hyperparameters of the model. It is important that the cross-validation process never sees the testing dataset in order to avoid information leakage and overfitting.

```
In [376]: #separate dataset into independent variables (X) and dependent variable (y)
X = full_data_balanced.drop(columns=['Status', 'Loan Title', 'State'])
y = full_data_balanced['Status']

#splitting the data
#choose a small training size to speed up the model fitting process
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, stratify = y)
print(X_train.shape,X_test.shape,y_train.shape,y_test.shape)

(476583, 63) (317723, 63) (476583,) (317723,)
```

2.3 Cross-validation

In the cross-validation process, only the training data is used. The LogisticRegressionCV function is used for cross-validation since it is more efficient than using the brute force parameter search provided by the GridSearchCV function.²

The LogisticRegressionCV function is mainly used to optimize for C, i.e. the inverse of regularization parameter values. When an integer is provided for the C parameter, e.g. Cs = 10, then a grid of Cs values are chosen in a logarithmic scale between 1e-4 and 1e4.³

I have chosen to perform cross-validation on:

- a logistic regression model using a liblinear solver with l1 penalty
- a logistic regression model using a liblinear solver with l2 penalty
- a logistic regression model using a lbfgs solver (lbfgs solver only supports l2 penalty, which is the default)

```
In [419]: from sklearn.linear_model import LogisticRegression,LogisticRegressionCV
from sklearn.metrics import classification_report

logr_lib_l1 = LogisticRegressionCV(Cs=10,solver='liblinear',penalty='l1').fit(X_train,
y_train)
logr_lib_l2 = LogisticRegressionCV(Cs=10,solver='liblinear',penalty='l2').fit(X_train,
y_train)
logr_lbfgs = LogisticRegressionCV(Cs=10,solver='lbfgs').fit(X_train, y_train)

In [420]: logr_models = [logr_lib_l1,logr_lib_l2,logr_lbfgs]

print('Logistic Regression CV results using different solvers, optimizing for C.\n')

for model in logr_models:
    print(f'Logistic Regression CV results using {str(model)}.\n')
    print(f'Array of C used for cross-validation: \n{model.Cs_}\n')
    print(f'Array of C that maps to the best scores across every class: \n{model.C_}\n')
    print(f'Grid of scores obtained during cross-validating each fold:
\n{model.scores_[1]}\n')
    print(f'Mean accuracy score: \n{np.mean(model.scores_[1])}\n')
    print(f'Classification report:\n{classification_report(y_train,
model.predict(X_train))}')

```

²https://scikit-learn.org/stable/modules/grid_search.html#multimetric-grid-search

³https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegressionCV.html#sklearn.linear_model.LogisticRegressionCV

Logistic Regression CV results using different solvers, optimizing for C.

```
Logistic Regression CV results using LogisticRegressionCV(Cs=10, class_weight=None,
cv=None, dual=False,
    fit_intercept=True, intercept_scaling=1.0, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l1', random_state=None,
    refit=True, scoring=None, solver='liblinear', tol=0.0001,
    verbose=0).
```

Array of C used for cross-validation:

```
[1.00000000e-04 7.74263683e-04 5.99484250e-03 4.64158883e-02
 3.59381366e-01 2.78255940e+00 2.15443469e+01 1.66810054e+02
 1.29154967e+03 1.00000000e+04]
```

Array of C that maps to the best scores across every class:

```
[2.7825594]
```

Grid of scores obtained during cross-validating each fold:

```
[[0.81104984 0.81417834 0.82075638 0.82226083 0.82255039 0.82251892
 0.82248115 0.82248744 0.82248744 0.82248744]
 [0.80908467 0.81206841 0.81825621 0.81957812 0.81961589 0.8197292
 0.81974179 0.81974808 0.81973549 0.81974808]
 [0.81011583 0.81342692 0.81963364 0.82093667 0.82104369 0.82106887
 0.82104998 0.82104369 0.82104369 0.82104369]]
```

Mean accuracy score:

```
0.819032358927715
```

Classification report:

	precision	recall	f1-score	support
0	0.78	0.90	0.83	238291
1	0.88	0.74	0.81	238292
avg / total	0.83	0.82	0.82	476583

```
Logistic Regression CV results using LogisticRegressionCV(Cs=10, class_weight=None,
cv=None, dual=False,
    fit_intercept=True, intercept_scaling=1.0, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
    refit=True, scoring=None, solver='liblinear', tol=0.0001,
    verbose=0).
```

Array of C used for cross-validation:

```
[1.00000000e-04 7.74263683e-04 5.99484250e-03 4.64158883e-02
 3.59381366e-01 2.78255940e+00 2.15443469e+01 1.66810054e+02
 1.29154967e+03 1.00000000e+04]
```

Array of C that maps to the best scores across every class:

```
[0.35938137]
```

Grid of scores obtained during cross-validating each fold:

```
[[0.81181151 0.81644446 0.8212096 0.82113407 0.82226712 0.82219159
 0.82225454 0.82226083 0.82226712 0.82227971]
 [0.81007296 0.81479407 0.81768968 0.8181492 0.81819956 0.81819956
 0.81819956 0.81934521 0.81820585 0.81821215]
 [0.81090268 0.81626589 0.81264006 0.82076671 0.82091149 0.82077301
 0.82076671 0.81962734 0.81972806 0.82076671]]
```

Mean accuracy score:
0.8186112343554276

Classification report:

	precision	recall	f1-score	support
0	0.77	0.90	0.83	238291
1	0.89	0.74	0.80	238292
avg / total	0.83	0.82	0.82	476583

Logistic Regression CV results using LogisticRegressionCV(Cs=10, class_weight=None, cv=None, dual=False, fit_intercept=True, intercept_scaling=1.0, max_iter=100, multi_class='ovr', n_jobs=1, penalty='l2', random_state=None, refit=True, scoring=None, solver='lbfgs', tol=0.0001, verbose=0).

Array of C used for cross-validation:
[[1.00000000e-04 7.74263683e-04 5.99484250e-03 4.64158883e-02
3.59381366e-01 2.78255940e+00 2.15443469e+01 1.66810054e+02
1.29154967e+03 1.00000000e+04]]

Array of C that maps to the best scores across every class:
[2.7825594]

Grid of scores obtained during cross-validating each fold:
[[0.8039934 0.8039934 0.8039934 0.8039934 0.8039934 0.8039997
0.8039934 0.8039934 0.8039934 0.8039934]
[0.80166938 0.80166938 0.80166938 0.80166938 0.80166938 0.80166938
0.80166938 0.80166938 0.80166938 0.80166938]
[0.80292711 0.80292711 0.80292711 0.80292711 0.80292711 0.80292711
0.80292711 0.80292711 0.80292711 0.80292711]]

Mean accuracy score:
0.802863507370328

Classification report:

	precision	recall	f1-score	support
0	0.76	0.90	0.82	238291
1	0.87	0.71	0.78	238292
avg / total	0.81	0.80	0.80	476583

2.4 Feature-scaling Models

I feature-scaled the independent variables to lie between 0 and 1 using a MinMaxScaler. Feature-scaling may improve the model if "noisy" variables had an unjustly high weight, while important variables had comparatively low weights. However, it may also decrease model accuracy if the "noisy" variables were already weighted low and the important variables were already weighted high in the original data; in such a case, feature-scaling would only emphasize the noise by giving it equal weightage as the important predictors.

```
In [413]: from sklearn import preprocessing
```

```
#try with feature scaling
min_max_scaler = preprocessing.MinMaxScaler()
X_train_scaled = min_max_scaler.fit_transform(X_train)

logr_lib_l1_scaled =
LogisticRegressionCV(Cs=10,solver='liblinear',penalty='l1').fit(X_train_scaled, y_train)
logr_lib_l2_scaled =
LogisticRegressionCV(Cs=10,solver='liblinear',penalty='l2').fit(X_train_scaled, y_train)
```

```
In [418]: logr_scaled_models = [logr_lib_l1_scaled,logr_lib_l2_scaled]
```

```
print('Logistic Regression CV results using liblinear solver with scaled features,
optimizing for C.\n')

for model in logr_scaled_models:
    print(f'Logistic Regression CV results using {str(model)}.\n')
    print(f'Array of C used for cross-validation: \n{model.Cs_}\n')
    print(f'Array of C that maps to the best scores across every class: \n{model.C_}\n')
    print(f'Grid of scores obtained during cross-validating each fold:
\n{model.scores_[1]}\n')
    print(f'Mean accuracy score: \n{np.mean(model.scores_[1])}\n')
    print(f'Classification report:\n{classification_report(y_train,
model.predict(X_train_scaled))}')

```

```
Logistic Regression CV results using liblinear solver with scaled features, optimizing
for C.
```

```
Logistic Regression CV results using LogisticRegressionCV(Cs=10, class_weight=None,
cv=None, dual=False,
    fit_intercept=True, intercept_scaling=1.0, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l1', random_state=None,
    refit=True, scoring=None, solver='liblinear', tol=0.0001,
    verbose=0).
```

```
Array of C used for cross-validation:
```

```
[1.00000000e-04 7.74263683e-04 5.99484250e-03 4.64158883e-02
 3.59381366e-01 2.78255940e+00 2.15443469e+01 1.66810054e+02
 1.29154967e+03 1.00000000e+04]
```

```
Array of C that maps to the best scores across every class:
```

```
[2.7825594]
```

```
Grid of scores obtained during cross-validating each fold:
```

```
[[0.78129446 0.81081064 0.81713689 0.82349461 0.82197757 0.82376528
 0.8231358 0.82256927 0.82250003 0.82248115]
 [0.77944241 0.80902802 0.81520323 0.82001246 0.82058529 0.82164282
 0.82015095 0.81974808 0.81975438 0.81974179]
 [0.78140501 0.81015359 0.81630366 0.82128919 0.82176759 0.82293214
 0.8218872 0.82113811 0.8210311 0.82104369]]
```

Mean accuracy score:
0.8157808796758563

Classification report:

	precision	recall	f1-score	support
0	0.78	0.91	0.84	238291
1	0.89	0.74	0.81	238292
avg / total	0.83	0.82	0.82	476583

Logistic Regression CV results using LogisticRegressionCV(Cs=10, class_weight=None, cv=None, dual=False, fit_intercept=True, intercept_scaling=1.0, max_iter=100, multi_class='ovr', n_jobs=1, penalty='l2', random_state=None, refit=True, scoring=None, solver='liblinear', tol=0.0001, verbose=0).

Array of C used for cross-validation:
[1.00000000e-04 7.74263683e-04 5.99484250e-03 4.64158883e-02
3.59381366e-01 2.78255940e+00 2.15443469e+01 1.66810054e+02
1.29154967e+03 1.00000000e+04]

Array of C that maps to the best scores across every class:
[10000.]

Grid of scores obtained during cross-validating each fold:
[[0.81154083 0.81285644 0.82007654 0.82100188 0.82124108 0.82137327
0.82202163 0.82245597 0.82317987 0.82300361]
[0.81008555 0.80911615 0.81792259 0.81952776 0.81982362 0.81996211
0.82040274 0.82065453 0.8208182 0.82097557]
[0.812596 0.8135717 0.82113811 0.82152209 0.82160393 0.82172983
0.82220823 0.82233413 0.82207604 0.82217046]]

Mean accuracy score:
0.8192996817342216

Classification report:

	precision	recall	f1-score	support
0	0.78	0.91	0.84	238291
1	0.89	0.74	0.81	238292
avg / total	0.83	0.82	0.82	476583

2.5 Best Hyperparameters

From the results, it appears that the logistic regression model using a liblinear solver with l2 penalty, with feature-scaling, has the highest accuracy, precision, and recall:

- accuracy: 0.819299
- precision: 0.83
- recall: 0.82

Without feature scaling, the logistic regression model using a liblinear solver with l1 penalty has the best performance:

- accuracy: 0.819032
- precision: 0.83
- recall: 0.82

Since the accuracy is only off by a negligible amount, and feature-scaling takes up additional resources, I will use the logistic regression model using a liblinear solver with l1 penalty to predict the results of the test data.

2.6 Model Evaluation

The chosen model is applied on the testing set to measure its performance. The accuracy, precision, and recall scores shown below suggest that the model is performing quite well.

```
In [428]: y_pred = logr_lib_l1.predict(X_test)
          print(f"Accuracy score: {logr_lib_l1.score(X_test,y_test)}\n")
          print(f"Classification report: \n{classification_report(y_test, y_pred)}")
```

Accuracy score: 0.8206236249815091

Classification report:

	precision	recall	f1-score	support
0	0.78	0.90	0.83	158862
1	0.88	0.74	0.80	158861
avg / total	0.83	0.82	0.82	317723

2.7 Predicting Highest Loan Amount for Approval

The predicted probabilities of being approved or rejected by the lending club based on the independent variables of amount requested, debt-to-income ratio, employment length, loan title (one-hot-encoded), and state (one-hot-encoded) can also be obtained from the logistic regression model. By setting a boundary of 0.95 approval as a safe threshold for approval, we can advise customers on the highest loan amount they could potentially borrow by inputting their data into the model, and tweaking the amount requested until the 0.95 approval boundary is achieved.⁴

The dataframe below displays a small snapshot of individual features that resulted in ≥ 0.95 probability of approval.

```
In [497]: y_proba_pred = logr_lib_l1.predict_proba(X_test)
X_test_df = pd.DataFrame(X_test).reset_index()
y_proba_pred_df = pd.DataFrame(y_proba_pred).reset_index()

predictions_df = pd.concat([X_test_df, y_proba_pred_df], axis = 1)
predictions_df.drop(columns=['index'], inplace=True)
predictions_df_high_approval = predictions_df[predictions_df[1] >= 0.95]
predictions_df_high_approval.head()
```

```
Out [497]:
```

	Amount Requested	Debt-To-Income Ratio	Employment Length	Business	Car	\
3	12000.0	12.0	10	0	0	
19	9000.0	18.0	10	0	0	
21	10000.0	14.0	10	0	0	
24	25000.0	4.0	10	0	0	
25	12000.0	4.0	10	0	0	

	Credit Card	Debt Consolidation	Green Loan	House	Major Purchase	\
3	0	0	0	0	1	
19	0	1	0	0	0	
21	0	1	0	0	0	
24	0	0	0	1	0	
25	0	0	0	0	0	

	...	TX	UT	VA	VT	WA	WI	WV	WY	0	1
3	...	0	0	0	0	0	0	0	0	0.037430	0.962570
19	...	0	0	1	0	0	0	0	0	0.016704	0.983296
21	...	0	0	0	0	0	0	0	0	0.013274	0.986726
24	...	0	0	0	0	0	0	0	0	0.024434	0.975566
25	...	0	0	0	0	1	0	0	0	0.039848	0.960152

[5 rows x 65 columns]

⁴**#probability:** I used the best logistic regression model out of all the trained models to predict the probability of loan approval on the unseen test set. The probability can be set at a specific threshold (in this case, 0.95), and the various features can then be manipulated to obtain the probability threshold that satisfies our decision boundary. If we had a higher tolerance for loan rejection, we could set the probability threshold to a lower level, e.g. 0.75. Probability is thus used as a way to gauge and predict the highest loan amount that we would safely advise to a customer.

2.7.1 Example Customer

The predictive model is tested on an example customer, Bob, who has a Debt-To-Income Ratio of 3, Employment Length of 7 years, is taking a loan for Debt Consolidation, and lives in California.

While testing the model, I experienced the following behavior:

- for a given set of customer stats, the model plateaus at a fixed probability of approval, insensitive to large amount changes.
- increasing the amount requested increased the probability of the loan being approved.

Some reasons for this behavior are conjectured as follows: - Other variables might play a much more significant role of predicting loan approval or rejection, and so changes in the amount requested do not do much to increase or decrease the customer's chances of getting a loan. - A higher amount requested could also correlate with better customer stats, and so the model might have ended up learning that with a given set of "approvable" stats, a higher requested amount means a higher rate of approval.

As opposed to changing the amount requested, a variable which seems to be extremely responsive in terms of affecting a customer's loan status is the employment length. A higher employment length starkly increases the probability of loan approval and vice versa. This makes intuitive sense since a person with longer employment length would most likely be able to pay off the loan since they have a constant, stable stream of income.

Both predictive behaviors for Bob are demonstrated below.

```
In [680]: bob = {
    'Amount Requested':10000,'Debt-To-Income Ratio':3,'Employment Length':7,
    'Business':0,'Car':0,'Credit Card':0,'Debt Consolidation':1,
    'Green Loan':0,'House':0,'Major
Purchase':0,'Medical':0,'Moving':0,'Other':0,'Vacation':0,
    'AK':0,'AL':0,'AR':0,'AZ':0,'CA':1,'CO':0,'CT':0,'DC':0,'DE':0,
    'FL':0,'GA':0,'HI':0,'IL':0,'IN':0,'KS':0,'KY':0,'LA':0,'MA':0,'MD':0,
    'ME':0,'MI':0,'MN':0,'MO':0,'MS':0,'MT':0,'NC':0,'ND':0,'NE':0,'NH':0,
    'NJ':0,'NM':0,'NV':0,'NY':0,'OH':0,'OK':0,'OR':0,'PA':0,'RI':0,'SC':0,
    'SD':0,'TN':0,'TX':0,'UT':0,'VA':0,'VT':0,'WA':0,'WI':0,'WV':0,'WY':0
}

bob_df = pd.DataFrame(bob, index=[0])

print("With Debt-To-Income Ratio of 3.0, Employment Length 7, Debt Consolidation and
CA:\n")
for i in range(10000,60000,10000):
    bob_df['Amount Requested'] = i
    bob_pred = logr_lib_l1.predict_proba(bob_df)
    print(f"Amount Requested: {i}. Probability of approval: {bob_pred[0][1]}\n")

print("With Amount Requested of 10000, Debt-To-Income Ratio of 3.0, Debt Consolidation
and CA:\n")
for i in range(0,11,2):
    bob_df['Employment Length'] = i
    bob_pred = logr_lib_l1.predict_proba(bob_df)
    print(f"Employment Length: {i}. Probability of approval: {bob_pred[0][1]}\n")
```

With Debt-To-Income Ratio of 3.0, Employment Length 7, Debt Consolidation and CA.:

Amount Requested: 10000. Probability of approval: 0.9432501811402284

Amount Requested: 20000. Probability of approval: 0.9444063842278166

Amount Requested: 30000. Probability of approval: 0.9455403912338444

Amount Requested: 40000. Probability of approval: 0.9466525732485206

Amount Requested: 50000. Probability of approval: 0.9477432973968143

With Amount Requested of 10000, Debt-To-Income Ratio of 3.0, Debt Consolidation and CA.:

Employment Length: 0. Probability of approval: 0.27943664927698353

Employment Length: 2. Probability of approval: 0.5377687181673515

Employment Length: 4. Probability of approval: 0.777297113653875

Employment Length: 6. Probability of approval: 0.9128232861071167

Employment Length: 8. Probability of approval: 0.9691483057945611

Employment Length: 10. Probability of approval: 0.9895002327730532

2.8 Data Visualizations - Feature values which lead to higher probability of approval

Below are some bar plots visualizing various feature values, ordered by the top 10 highest mean probabilities of approval. While most of the value changes don't appear to lead to significant changes in the probability, there are some categories - most noticeably "Employment Length" - that display a relatively clear relationship between the feature value and the mean probability of being approved. For "Employment Length", we can see that as the "Employment Length" decreases, the probability of being accepted for a loan also decreases.

Only a few feature categories were plotted since this is more for interest and observation vs a rigorous form of analysis; the multidimensional nature of the problem means that plotting single features against the probability will not provide robust results. Also, do note that the x-axis is not in order, and the y-axis has a trimmed scale in order to show the minute differences between the data. For best results, the trained model should be used directly.

```
In [581]: import seaborn as sns
          sns.set()

          plt.figure(figsize=(20,30))

          plt.subplot(4,2,1)
          predictions_df[1].groupby(predictions_df['Amount
          Requested']).mean().nlargest(10).plot.bar(ylim=(0.96,1))

          plt.subplot(4,2,2)
          predictions_df[1].groupby(predictions_df['Debt-To-Income
          Ratio']).mean().nlargest(10).plot.bar(ylim=(0.59,0.61))

          plt.subplot(4,2,3)
          predictions_df[1].groupby(predictions_df['Employment
          Length']).mean().nlargest(10).plot.bar(ylim=(0,1))

          plt.subplot(4,2,4)
```

```


predictions_df[1].groupby(predictions_df['Business']).mean().nlargest(10).plot.bar(ylim=(0,1))

plt.subplot(4,2,5)
predictions_df[1].groupby(predictions_df['Green
Loan']).mean().nlargest(10).plot.bar(ylim=(0,1))

plt.subplot(4,2,6)
predictions_df[1].groupby(predictions_df['House']).mean().nlargest(10).plot.bar(ylim=(0,
1))

```

Out[581]: <matplotlib.axes._subplots.AxesSubplot at 0x1571b1668>



3 Summary

In this report, I built a logistic regression model that can predict the probability of a loan being approved or rejected based on various features provided by the customer.

3.1 Variables included in the model

The independent variables included in the model were: Amount Requested, Debt-To-Income Ratio, Employment Length, Loan Title (one-hot-encoded), State (one-hot-encoded). The response variables were the status of the loan, with 1 signifying an approved loan and 0 signifying a rejected loan.

3.2 Cleaning and transformations on the data

The following steps were taken to clean and transform the data for training and testing:

- Dropping columns that were not present in both the approval and rejection data.
- Dropping columns with a high amount of irregular values (more than half).
- Dropping rows with NaNs.
- Converting variables to the appropriate data types.
- Dropping categorical variables with a high number of discrete categories (in this case, 998 categories for unique zip codes).
- One-hot-encoding categorical variables that were kept (loan titles, states).
- Balancing the response variable class so that the ratio of approval to rejections is 1.0 in the entire dataset, and the same balanced ratio is transferred into the training and test sets.
- Feature scaling the independent variables to see whether that improves the model's performance.

3.3 The type of model used and model settings

A logistic regression model was used to predict the probability of getting a loan approval since logistic regression is suitable for modeling binary dependent variables. Cross-validation was conducted to finetune the best hyperparameters for the model. The parameters tuned were the Cs, i.e. the inverse of regularization parameter values; the normalization penalty, i.e. L1 or L2; and the type of solver, e.g. liblinear, lbfgs (limited memory Broyden-Fletcher-Goldfarb-Shann algorithm). The final model settings chosen - based on the model's performance on the training data - are C: 2.78, solver: liblinear, penalty: L1.

With feature-scaled variables, the best performing model had the following parameters: C: 10000, solver: liblinear, penalty: L2. The two best performing models (with feature-scaling vs without feature-scaling) only differed by a miniscule amount in terms of accuracy score (0.81929 vs 0.81903 respectively), and had the same average precision and recall scores of 0.83 and 0.82. Since feature-scaling would require additional computational resource, I have decided to use the model that did not require feature-scaling (C: 2.78, liblinear, L1 penalty).⁵

⁵**#regression:** The bulk of this assignment hinges on the development and training of a logistic regression model to predict the probability of loan approval for a given customer based on various independent variables. I provided motivation for the use of logistic regression (modeling binary dependent variables), used cross-validation to finetune various hyperparameters of the model, and trained the model on both feature-scaled and non-feature-scaled data to obtain the model with the best performance to be used on unseen test data. The chosen model is then used to predict

3.4 Training method used, and techniques to avoid overfitting the data

To train the model, the data was split into a training and test set with 0.4 test size and 0.6 training size, i.e. 317723 data points for the test set, and 476583 for the training set. By only performing cross-validation on the training set, overfitting is prevented since the test set remains unseen by the model until the final performance testing stage. The LogisticRegressionCV from the sklearn library is used for cross-validation since it is more efficient than using a brute force search over the parameter space, e.g. GridSearchCV.⁶ The LogisticRegressionCV function is used with different solvers and optimizes for C. The best performing model is then chosen to predict results for the test data.

3.5 Estimate of how well the model will perform on unseen data

The model has a relatively good performance on the test data, returning an accuracy of 0.82, average precision of 0.83, and average recall of 0.82. As such, I infer that the model would show a similar performance on a different set of unseen data.

3.6 Assumptions

- The model was only trained on 2015 data, so the performance might be subpar when used on data with a stark difference in time frame.
- The model might be even more accurate with less features (some of these features could just be noise), but this was not tested.
- Some significant features may have been left out during the data cleaning process. As such, it is difficult to determine whether there were confounding features that could be affecting the data. That being said, the model does appear to work well as a predictive model.
- The loan titles were manually categorized by my personal judgement. While I believe that I have categorized them reasonably, I could be unaware that some categories should be more appropriately categorized together/separated.

the probability of approval for a manually constructed customer.

⁶https://scikit-learn.org/stable/modules/grid_search.html#grid-search

4 Appendix

Code in Jupyter Notebook format is available online at:

<https://github.com/hueyning/cs156-ml/blob/master/assignment-2-lending-data/assignment-2-lending-data.ipynb>